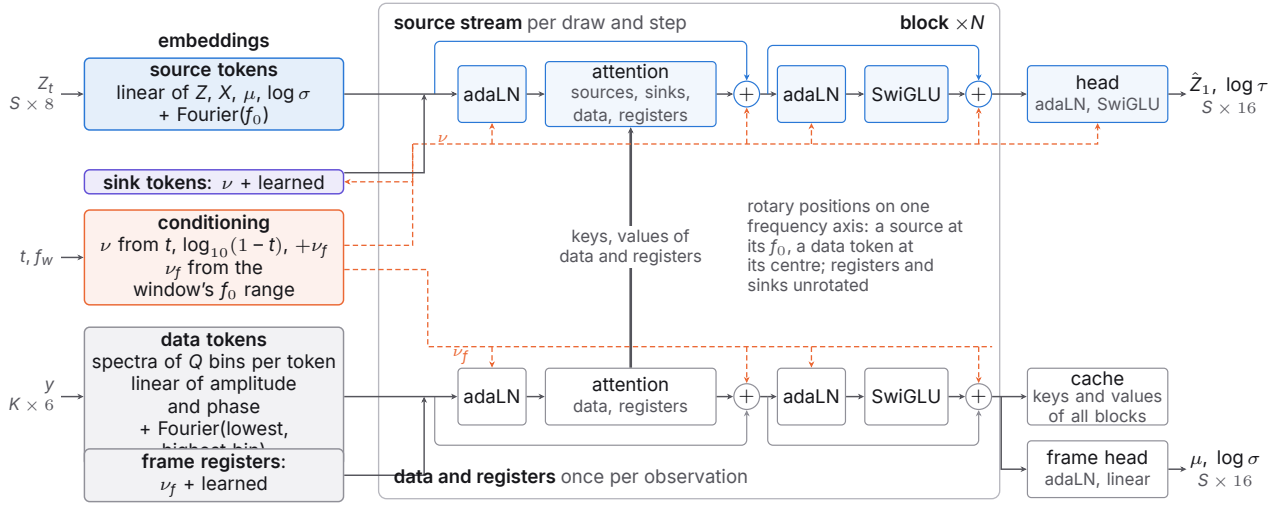
embeddings
source stream per draw and step
block ×N
Z_t
$S \times 8$
source tokens
linear of Z, X, μ, log σ
+ Fourier(f_0)
adaLN
attention
sources, sinks, data, registers
adaLN
SwiGLU
head
adaLN, SwiGLU
$\hat{Z}_1$, log τ
$S \times 16$
ν
sink tokens: ν + learned
conditioning
ν from t, $\log_{10}(1-t)$, $+\nu_f$
ν_f from the window's f_0 range
t, f_w
keys, values of data and registers
rotary positions on one frequency axis: a source at its f_0, a data token at its centre; registers and sinks unrotated
ν_f
y
$K \times 6$
data tokens
spectra of Q bins per token
linear of amplitude and phase
+ Fourier(lowest,
adaLN
attention
data, registers
adaLN
SwiGLU
cache
keys and values of all blocks
frame registers:
ν_f + learned
data and registers once per observation
frame head
adaLN, linear
μ, log σ
$S \times 16$