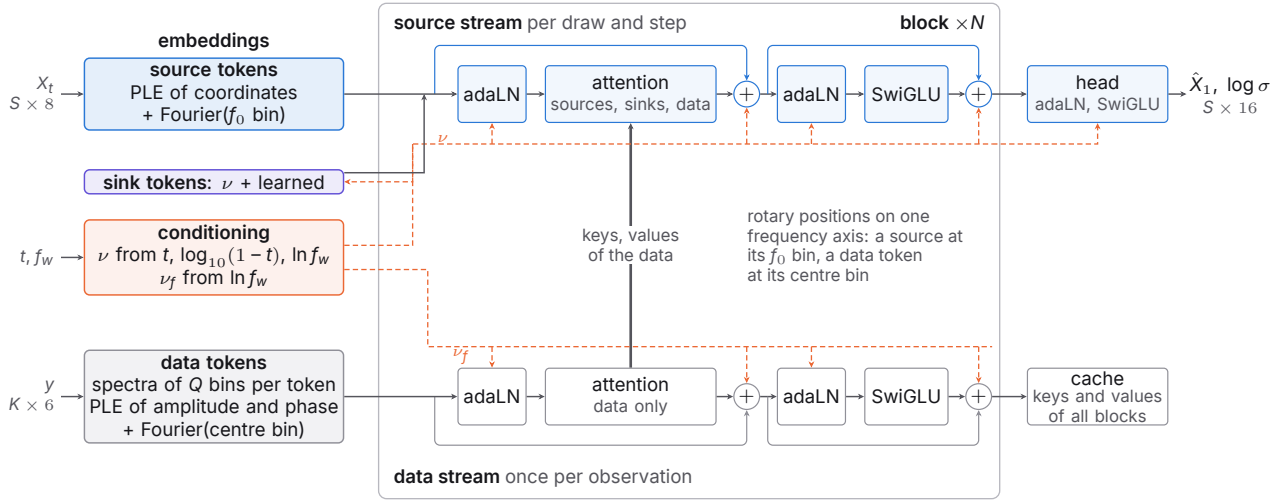
embeddings
source stream per draw and step
block ×N
X_t
$S \times 8$
source tokens
PLE of coordinates
+ Fourier(f_0 bin)
adaLN
attention
sources, sinks, data
adaLN
SwiGLU
head
adaLN, SwiGLU
$\hat{X}_1$, $\log \sigma$
$S \times 16$
ν
sink tokens: ν + learned
conditioning
ν from t, $\log_{10}(1-t)$, $\ln f_w$
ν_f from $\ln f_w$
t, f_w
keys, values
of the data
rotary positions on one
frequency axis: a source at
its f_0 bin, a data token
at its centre bin
ν_f
y
$K \times 6$
data tokens
spectra of Q bins per token
PLE of amplitude and phase
+ Fourier(centre bin)
adaLN
attention
data only
adaLN
SwiGLU
cache
keys and values
of all blocks
data stream once per observation